

DAN MATHIESON

Claude Code Platform Builder | Technical Enablement | AI Training & Adoption

danny.mathieson233@gmail.com | danielmathieson.com | linkedin.com/in/daniel-mathieson | San Francisco, CA

May 18, 2026

Claude Code Enablement Team

Anthropic

San Francisco, CA

Dear Anthropic Hiring Team,

The most underrated moment in any AI adoption curve is the one where someone who is not an AI researcher decides they trust a system enough to let it change how they work. Getting to that moment is not about having a better demo. It is about meeting people where they actually are: showing the capability in the context of the specific workflow they run every day, giving them a way to understand what the system is actually doing, and building the kind of iterative feedback loop that makes the tool better over time as it learns from real use. I have spent two years building exactly that, inside a production GTM org running on Claude Code. I want to do it at Anthropic's scale for developer and engineering teams.

At Smarter Technologies, I built a **22-skill, 5-agent production platform on Claude Code** from scratch. The design principle from day one was that the enablement infrastructure and the platform itself were the same artifact: every skill had a YAML manifest that served simultaneously as configuration spec and onboarding documentation. I implemented pre- and post-tool-use hooks to sample feedback on **15-20% of skill invocations**, achieving a **93% tool approval rating** in the first three weeks of cross-company deployment. The feedback signal (clusters of thumbs-down on the pricing skill with notes about wrong EHR context) told me exactly what to fix in the next iteration. I then led the deployment of a unified Cowork platform across all operational teams, sales, GTM, and delivery, using Anthropic's GitHub-based plugin repository to distribute role-specific skill sets with no prompt-writing required. That feedback-loop-as-product discipline is something I would bring to building and iterating on Claude Code enablement content at scale.

The pattern runs through my entire career, not just the last two years. At FanDuel I designed and taught SQL courses to the operations team, writing curriculum from scratch to match the specific reporting workflows they ran every day. At Google I designed and delivered Apps Script training to the finance team alongside the automation infrastructure itself, so that when the system was handed over, the team could maintain and extend it independently. When I hired SE candidates at Smarter Technologies, I built demo coaching rubrics to evaluate technical communication against real prospect scenarios. When I deployed Cowork to non-technical sales reps, I wrote skill specifications that assumed no prompt-writing experience and built onboarding docs that walked users through their first three use cases with concrete examples. Every piece of technical work I have done has had a "how do I make this learnable for the people who need it" layer on top of it.

At Thoughtful AI, I built the company's first customer-facing analytics product on AWS QuickSight, parsing agent execution logs into dashboards used in executive QBRs. Once the CEO and CPO saw how those dashboards changed what was possible in an enterprise sales conversation, they started bringing me into every large deal.

96% of prospects who met with me moved to proposal within a month, not because I had a better pitch, but because I could connect the technology to the specific workflow the customer was actually trying to fix. I was promoted to Director of Sales Engineering the day the sales team asked to have me full-time. The lesson from that arc: making complex AI systems legible is not a soft skill. It is the thing that determines whether the technology gets deployed or not.

My personal site runs a production agent built on the Anthropic SDK. It leverages tool use, streaming, and a RAG pipeline grounded in copious amounts of career content I've written. This is the same kind of Claude-native architecture I would be helping enterprise customers design. It is at [danielmathieson.com](https://www.danielmathieson.com) if you want to see it before we talk! For a review of all the roles I've considered at Anthropic, you can navigate to <https://www.danielmathieson.com/refer/f3a8c1d94e72>.

When Smarter Technologies shut down, I was retained as one of the 15% kept, the only non-Product/Engineering person, specifically because I can take complex AI systems and make them legible to people who need to adopt them. My enablement work so far has been internal to GTM orgs rather than external developer curriculum delivery. That is a genuine gap I would be transparent about. But the underlying teaching instinct and the production Claude Code depth are the hard parts, and I have both. I would welcome the chance to show that in a conversation.

Sincerely,

Dan Mathieson

danny.mathieson233@gmail.com · [danielmathieson.com](https://www.danielmathieson.com)