

# DAN MATHIESON

Claude Code Platform Builder | Enterprise AI Enablement | Technical Solutions Engineering

danny.mathieson233@gmail.com | danielmathieson.com | linkedin.com/in/daniel-mathieson | San Francisco, CA

---

May 18, 2026

## Hiring Team

Claude Code, Go-To-Market

Anthropic

San Francisco, CA

Dear Anthropic Hiring Team,

The hardest part of enterprise AI adoption is not the technology. It is closing the gap between what the platform can actually do and what the people responsible for deploying it are willing to trust. The best way to earn that trust is not a better demo. It is honesty: being specific about where AI genuinely helps, direct about where it does not fit a given workflow, and grounding every conversation in the customer's actual business context rather than the product's capabilities in the abstract. That mix shifts over time as the technology improves, and part of the job is continuously reevaluating it with customers. I have spent the past two years doing exactly that: building production AI systems and then teaching non-technical enterprise teams to trust and use them. I want to do that work at Anthropic's scale.

At Smarter Technologies, I built a production platform on Claude Code from scratch, using it as the primary operating system for a \$31.8M enterprise pipeline. I implemented pre- and post-tool-use hooks to sample user feedback on 15-20% of skill invocations and feed that signal directly back into skill iteration, achieving a 93% tool use approval rating across sampled interactions. Then I led the initiative to extend this beyond my own team: I designed and deployed a unified Cowork-based plugin platform covering all operational teams across Thoughtful and Smarter Technologies, sales, GTM, and delivery, consolidating independently-built Claude Code work across the org into one governed package using Anthropic's GitHub-based plugin repository system to distribute role-specific skill sets with no prompt-writing required. I also built an event-driven agent that waited for the engineering team's weekly standup pages to be created in Notion and auto-populated a market feedback section from that week's discovery calls, closing the loop between field conversations and engineering priorities automatically. When team members started Slacking me to ask my agent questions on their behalf, I knew the adoption barrier was real and solvable. Building that solution is what this role does, except at the scale of hundreds of enterprise customers instead of one company.

My path to building that platform started at Thoughtful AI, where I developed the skill this role depends on most: making complex AI legible to the people who need to adopt it. I built Thoughtful's first customer-facing analytics product on AWS QuickSight, parsing agent execution logs into dashboards used in executive QBRs. That product became the core GTM sales artifact showing what separated our agent observability from generic RPA. Once the CEO and CPO saw how those dashboards landed in customer conversations, they started bringing me into every large sales opportunity. Within a few weeks I was on all of them. 96% of prospects who met with me moved to proposal stage within a month, not because I had a better pitch, but because I could connect the technology to the specific workflow the customer was actually trying to fix. I was promoted to Director of Sales Engineering the day the sales team asked to have me full-time.

My personal site runs a production agent built on the Anthropic SDK. It leverages tool use, streaming, and a RAG pipeline grounded in copious amounts of career content I've written. This is the same kind of Claude-native

architecture I would be helping enterprise customers design. It is at [danielmathieson.com](https://www.danielmathieson.com) if you want to see it before we talk! For a review of all the roles I've considered at Anthropic, you can navigate to <https://www.danielmathieson.com/refer/f3a8c1d94e72>.

That design discipline applies equally to the human side of every system I build. Anthropic has done exceptional work at two levels: making AI genuinely powerful for developers through Claude Code, and making it accessible to people broadly through Claude and Claude Cowork. Where I have found my edge is going a layer or two deeper than either. A prior authorization coordinator who spends eight hours a day on that one workflow needs AI to do something very specific, and the answer is different from what a denial management specialist needs, which is different again from what a billing team lead needs. Generic AI implementations rarely survive contact with that specificity. The gap between a team that adopts AI and a team that stops using it after a month is almost always whether the tool was designed with their particular job function in mind or designed for a broad user and asked to serve a specific one. Understanding those small pockets of end users deeply enough to close that gap is what I have been building toward, and it is the thing I most want to bring to Anthropic's enterprise customers.

I am based in San Francisco and actively looking forward to the travel this role requires. Going on-site is often the thing that actually moves a customer from skeptical to committed, and I have never minded adjusting my schedule to make that happen. At Thoughtful AI I led forward-deployed engineering teams based in South America, South Africa, and Eastern Europe, coordinating delivery against customer commitments across time zones regardless of what hour that required. At AccessHealthcare I now work daily with India-based delivery teams. The job description asks for someone who can build reference architectures, run technical workshops, close the field-to-product feedback loop, and help enterprise engineering teams move from "we tried it" to "we run it daily." Those are things I have been doing. I want to do them for Claude Code. I would welcome the chance to talk.

Sincerely,

**Dan Mathieson**

[danny.mathieson233@gmail.com](mailto:danny.mathieson233@gmail.com) · [danielmathieson.com](https://www.danielmathieson.com)